

Yuni Wu

South San Francisco, CA

yunihsuanwu@gmail.com | (720) 487-5403 | linkedin/wih | github/yuni-wyx | portfolio

Summary

AI Systems Engineer focused on making AI product behavior observable through evaluation, tool routing, streaming responses, and debugging signals. Builds traceable workflows that turn model outputs into reliable product behavior before users depend on them.

Experience

Graduate Research Assistant

Sept 2025 – May 2026

CAIRO Lab, University of Colorado Boulder — Embodied AI and robotics research lab

- Built scenario-level workflows comparing VLM navigation decisions for consistency and deployment robustness.
- Benchmarked embodied-AI navigation diagnostics across structured human-interaction scenarios.

Machine Learning Engineer Intern

May 2025 – Aug 2025

MyEdMaster LLC — AI healthcare and education startup

- Built repeatable model-quality evaluation workflows for a PyTorch/torchvision ResNet classifier on 50K+ samples, achieving 0.92 F1-score in one setup.
- Instrumented class-imbalance reliability checks with cross-validation, ROC-AUC, calibration, and regression validation.
- Mapped asynchronous Cloud Run inference paths to identify deployment, retry, and observability tradeoffs.

Software Engineer

June 2021 – Aug 2022

Shopee Co. Ltd. — Leading Southeast Asian e-commerce platform

- Reduced QA workload by 80% by engineering SQL validation and anomaly-detection pipelines on 800+ QA scenarios.
- Accelerated reporting speed by 78% with regression monitoring pipelines that surfaced data inconsistencies earlier.
- Automated API/Selenium data collection workflows to support validation, monitoring, and reliability analysis.

Projects

Stock Intelligence Copilot — Traceable AI Workflow Platform

[demo]

- Routed intent classification, tool selection, evidence aggregation, and response generation through LangChain/LCEL chains.
- Improved traceability, debugging, and observability with FastAPI SSE streams and structured execution schemas.
- Unified market data, fundamentals, news, earnings, and portfolio tools into evidence bundles with tool-level error capture.

EmotionFlow — LLM Recommendation Evaluation System

[demo]

- Evaluated real-time recommendation quality with OpenAI API-backed pipelines and prompt optimization workflows.
- Reduced hallucination by 34% using feedback-driven loops with preference signals and structured response analysis.
- Increased recommendation acceptance to 57% through A/B testing and iterative prompt optimization.

NailAI — Cloud-Native Inference Workflow

[demo]

- Enabled asynchronous image processing through Cloud Run services, Pub/Sub routing, and idempotent retries.
- Achieved 0% failure at 200 concurrent users with p95 latency of 6.8s through fault-tolerant orchestration.

Technical Skills

Languages: Python, SQL, Go, TypeScript

AI Systems: OpenAI API, LangChain/LCEL, LLM API integration, tool routing, response synthesis, evidence aggregation

Evaluation: prompt optimization, response analysis, feedback-driven evaluation loops, A/B testing, model evaluation

Backend & Cloud: FastAPI, REST APIs, Server-Sent Events (SSE), Cloud Run, Pub/Sub, BigQuery, Docker

Reliability: observability, structured logging, regression validation, automated testing, fallback-safe API workflows

Education

University of Colorado Boulder — MS, Data Science | GPA: 3.76/4.0

May 2026

Soochow University — BS, Data Science | GPA: 3.7/4.0

June 2021